

Ayers AskDocs Phase 5 Summary Table

Input: data/ayers_askdocs/phase3_chatgpt_expert_scored.jsonl | Reward: healthcare-professional quality/empathy mean | Bootstrap: 100 | Seed: 0
Behavior policy: physician response | Target policy: logged ChatGPT response | Human-rated ground truth, no LLM judge scoring.

Source	n	Config	Estimator	V_b	V_agent	Effect	V_hat	RelBias	RMSE	95% CI	Dir%	Cov
ayers_askdocs_2023	195	BGE-M3	DM	0.4254	0.7233	+0.2979	0.5537	-23.45%	0.1832	[0.502, 0.593]	100.00%	no
ayers_askdocs_2023	195	BGE-M3	MIPS	0.4254	0.7233	+0.2979	0.6193	-14.38%	0.1134	[0.555, 0.696]	100.00%	no
ayers_askdocs_2023	195	BGE-M3	OffCEM	0.4254	0.7233	+0.2979	0.5551	-23.25%	0.1826	[0.503, 0.593]	100.00%	no
ayers_askdocs_2023	195	BGE-M3 + learned concat-PCA	DM	0.4254	0.7233	+0.2979	0.5593	-22.67%	0.1611	[0.531, 0.597]	100.00%	no
ayers_askdocs_2023	195	BGE-M3 + learned concat-PCA	MIPS	0.4254	0.7233	+0.2979	0.6464	-10.63%	0.0857	[0.588, 0.716]	100.00%	no
ayers_askdocs_2023	195	BGE-M3 + learned concat-PCA	OffCEM	0.4254	0.7233	+0.2979	0.5607	-22.48%	0.1603	[0.532, 0.597]	100.00%	no

Notes

- Best by RMSE: BGE-M3 + learned concat-PCA + MIPS, RMSE 0.0857, V_hat 0.6464.
- Plain BGE-M3 MIPS: V_hat 0.6193, bias -0.1040, RMSE 0.1134.
- Learned concat-PCA MIPS: V_hat 0.6464, bias -0.0769, RMSE 0.0857.
- Both configurations recover the positive direction on all bootstrap resamples; neither plain bootstrap CI covers V_agent.