

Ayers AskDocs Quality-Only Phase 5 Summary Table

Input: data/ayers_askdocs_quality/phase3_chatgpt_expert_scored.jsonl | Reward: healthcare-professional quality mean only | Bootstrap: 100 | Seed: 0
Behavior policy: physician response | Target policy: logged ChatGPT response | Human-rated ground truth, no LLM judge scoring.

Source	n	Config	Estimator	V_b	V_agent	Effect	V_hat	RelBias	RMSE	95% CI	Dir%	Cov
ayers_askdocs_2023	195	BGE-M3	DM	0.5641	0.7829	+0.2188	0.6914	-11.68%	0.1102	[0.635, 0.720]	100.00%	no
ayers_askdocs_2023	195	BGE-M3	MIPS	0.5641	0.7829	+0.2188	0.7304	-6.71%	0.0612	[0.679, 0.776]	100.00%	no
ayers_askdocs_2023	195	BGE-M3	OffCEM	0.5641	0.7829	+0.2188	0.6926	-11.53%	0.1098	[0.635, 0.720]	100.00%	no
ayers_askdocs_2023	195	BGE-M3 + learned concat-PCA	DM	0.5641	0.7829	+0.2188	0.6983	-10.80%	0.0857	[0.668, 0.727]	100.00%	no
ayers_askdocs_2023	195	BGE-M3 + learned concat-PCA	MIPS	0.5641	0.7829	+0.2188	0.7587	-3.09%	0.0363	[0.706, 0.801]	100.00%	yes
ayers_askdocs_2023	195	BGE-M3 + learned concat-PCA	OffCEM	0.5641	0.7829	+0.2188	0.6998	-10.61%	0.0848	[0.669, 0.727]	100.00%	no

Notes

- Best by RMSE: BGE-M3 + learned concat-PCA + MIPS, RMSE 0.0363, V_hat 0.7587.
- Plain BGE-M3 MIPS: V_hat 0.7304, bias -0.0526, RMSE 0.0612.
- Learned concat-PCA MIPS: V_hat 0.7587, bias -0.0242, RMSE 0.0363; CI covers V_agent.
- Quality-only ground truth: V_b=0.5641, V_agent=0.7829, effect=+0.2188.